

cleaning_strings

October 19, 2020

Cleaning String Data in R

1 1. What is Data Cleaning?

- No uniform definition “data cleaning”
- Roughly speaking, refers to exploring the idiosyncrasies of a data set, and then addressing them in a principled manner so as to allow for data analysis

1.1 1.1 Examples of Data Cleaning

- Recoding "NULL", " ", "", to be NA
- Eliminating duplicate entries
- Ensure numeric data is being treated as numerics (e.g., "2" + 2 != 4)
- Treating dates or timestamps as Date or POSIXct data type

2 2. Cleaning Strings

- Parsing/cleaning/extracting info from strings is extremely common
- Parsing timestamp strings is a great example

2.1 2.1 Errors in our officer_cnt

```
In [1]: # Load necessary packages and arrests data
library(stringr)
library(dplyr)

arrests_df <- read.csv("./data/pvd_arrests_2020-10-03.csv")
```

Attaching package: dplyr

The following objects are masked from package:stats:

filter, lag

The following objects are masked from package:base:

```
intersect, setdiff, setequal, union
```

```
In [2]: count_names <- function(names_str) {  
  # This function should return the number of names in  
  # the string `names_str` that we pass to the function.  
  
  name_vec <- unlist(str_split(names_str, ", "))  
  k <- length(name_vec)  
  
  return(k)  
}
```

2.1.1 2.1.1 Inconsistencies in arresting_officers Column

```
In [3]: head(arrests_df$arresting_officers, 10)
```

1. 'YGonzalez, LTaveras' 2. 'NManfredi' 3. 'MPlace, JPerez, ASantos' 4. 'MPlace, JPerez, ASantos' 5. 'MPlace, JPerez, ASantos' 6. 'MPlace, JPerez, ASantos' 7. 'MPlace, JPerez, ASantos' 8. 'CVingi, SCooney' 9. 'CVingi, SCooney' 10. 'CVingi, SCooney'

```
In [4]: tail(arrests_df$arresting_officers, 10)
```

1. 'Lopez, Vincent/ Schneider, Alex/ Vargas, Guillermo' 2. 'Lopez, Vincent/ Schneider, Alex/ Vargas, Guillermo' 3. 'Lopez, Vincent/ Schneider, Alex/ Vargas, Guillermo' 4. 'Lopez, Vincent/ Schneider, Alex/ Vargas, Guillermo' 5. 'Maycock, Michael' 6. 'San Lucas, Luis' 7. 'Lopes, Joseph' 8. 'Lopes, Joseph' 9. 'Heaton, Robert' 10. 'Newton, Frank/ Chin, Rosemarie'

2.2 2.2 Addressing the Inconsistency

- Use different criteria for counting names with full-name format
 - Define function to identify full-name vs. first-initial format
 - Note: first-initial format always starts with two capital letters

```
In [5]: LETTERS # This is a built-in object in R
```

1. 'A' 2. 'B' 3. 'C' 4. 'D' 5. 'E' 6. 'F' 7. 'G' 8. 'H' 9. 'I' 10. 'J' 11. 'K' 12. 'L' 13. 'M' 14. 'N' 15. 'O' 16. 'P' 17. 'Q' 18. 'R' 19. 'S' 20. 'T' 21. 'U' 22. 'V' 23. 'W' 24. 'X' 25. 'Y' 26. 'Z'

```
In [6]: "B" %in% LETTERS
```

TRUE

2.3 2.3 Identifying Full-Name Format

- If the first two characters are uppercase, it's full-name format

```
In [7]: is_uppercase <- function(chr) {  
      res <- chr %in% LETTERS  
      return(res)  
    }  
  
    has_full_names <- function(names_str) {  
      char1 <- substr(names_str, 1, 1)  
      char2 <- substr(names_str, 2, 2)  
  
      res <- !(is_uppercase(char1) && is_uppercase(char2))  
      return(res)  
    }
```

2.3.1 2.3.1 Testing our Functions

```
In [8]: is_uppercase("a")           # false  
      is_uppercase("b")           # false  
      has_full_names("NManfredi")  # Not full name  
      has_full_names("MPlace, JPerez, ASantos") # Not full name  
  
      is_uppercase("A")  
      is_uppercase("B")  
      has_full_names("Newton, Frank")  
      has_full_names("Newton, Frank/ Chin, Rosemarie")  
  
FALSE  
FALSE  
FALSE  
FALSE  
TRUE  
TRUE  
TRUE  
TRUE
```

2.4 2.4 Fixing our count_names() Function

```
In [9]: old_count_names <- function(names_str) {  
      name_vec <- unlist(str_split(names_str, ", "))  
      k <- length(name_vec)  
  
      return(k)  
    }  
  
In [10]: count_names <- function(names_str) {  
      names_str_trm <- str_trim(names_str) # remove whitespace
```

```

    if (has_full_names(names_str_trm)) {
      split_char <- "/"
    } else {
      split_char <- ", "
    }

    name_vec <- unlist(str_split(names_str_trm, split_char))
    k <- length(name_vec)

    return(k)
  }

```

2.4.1 2.4.1 Testing New count_names()

```

In [11]: old_count_names("YGonzalez, LTaveras") == 2
old_count_names("Newton, Frank/ Chin, Rosemarie") == 2      # function is wrong
count_names("YGonzalez, LTaveras") == 2
count_names("Newton, Frank/ Chin, Rosemarie") == 2

TRUE
FALSE
TRUE
TRUE

```

2.5 2.5 Re-Counting Officers

- Let's compare how the "old" (i.e., incorrect) method did relative to our new count_names()

```

In [12]: count_officers <- function(col, old = FALSE) {

  n <- length(col)      # get the length of our input column
  cnts <- rep(0, n)     # allocate vector of zeros to populate with counts

  for (i in 1:n) {
    if (old) {
      cnts[i] <- old_count_names(col[i])
    } else {
      cnts[i] <- count_names(col[i])
    }
  }
  return(cnts)
}

```

```

In [13]: arrests_df$old_officer_cnt <- count_officers(arrests_df$arresting_officers, old = TRUE)

```

```

arrests_df$officer_cnt <- count_officers(arrests_df$arresting_officers)

```

```

In [14]: head(arrests_df)

```

		arrest_date <chr>	year <int>	month <int>	gender <chr>	race <chr>	ethnicity <chr>	year_of <int>
A data.frame: 6 × 20	1	2019-08-24T02:23:00.0	2019	8	Male	White	NonHispanic	1981
	2	2019-08-24T02:02:00.0	2019	8				1994
	3	2019-08-24T02:02:00.0	2019	8	Female	Black	NonHispanic	1984
	4	2019-08-24T02:02:00.0	2019	8	Female	Black	NonHispanic	1984
	5	2019-08-24T02:02:00.0	2019	8	Female	Black	Unknown	2001
	6	2019-08-24T02:02:00.0	2019	8	Female	Black	Unknown	2001

In [15]: `tail(arrests_df, 12)`

		arrest_date <chr>	year <int>	month <int>	gender <chr>	race <chr>	ethnicity <chr>	ye <in
A data.frame: 12 × 20	8744	2020-09-25T15:14:00.0	2020	9	Male	White	NonHispanic	19
	8745	2020-09-25T14:36:00.0	2020	9	Male	White	Hispanic	19
	8746	2020-09-25T14:36:00.0	2020	9	Male	White	Hispanic	19
	8747	2020-09-25T14:36:00.0	2020	9	Male	White	Hispanic	19
	8748	2020-09-25T14:36:00.0	2020	9	Male	White	Hispanic	19
	8749	2020-09-25T14:36:00.0	2020	9	Male	White	Hispanic	19
	8750	2020-09-25T09:45:00.0	2020	9	Male	White	NonHispanic	19
	8751	2020-09-25T09:11:00.0	2020	9	Male	Black	NonHispanic	19
	8752	2020-09-25T00:00:00.0	2020	9	Female	Black	Hispanic	19
	8753	2020-09-25T00:00:00.0	2020	9	Male	Black	Hispanic	19
	8754	2020-09-12T20:03:00.0	2020	9	Male	NULL	NULL	19
	8755	2020-08-27T07:10:00.0	2020	8	Male	White	NonHispanic	19

2.6 2.6 How Many Errors?

In [16]: `sum(arrests_df$old_officer_cnt != arrests_df$officer_cnt)`

4197

In [17]: `nrow(arrests_df)`

8755