

# intro\_ggplot2

October 19, 2020

Introduction to ggplot2

## 1 1. What is *ggplot2*?

- Hugely popular R package for visualization
- Authored by Hadley Wickham (of *dplyr* and *tidyverse* fame)
- Implements the “grammar-of-graphics” design philosophy (hence “gg”)
- Easily produces beautiful and informative visualizations

## 2 2. Plotting Person-Level Characteristics in Arrests

- The Pvd arrests data are at the *violation*-level
- We want person-level data on the individuals arrested
- This “*level-of-analysis*” or “*level-of-granularity*” problem is ubiquitous

## 3 3. Generating Person-Level Data

- We are aggregating “up” from the violation level
- Will use the `group_by()` and `summarise()` idiom

```
In [2]: # Load necessary packages
library(stringr)
library(dplyr)
library(ggplot2)

arrests_df <- read.csv("../data/pvd_arrests_2020-10-03.csv")
```

### 3.1 3.1 Computing Number of Officers (correctly)

- First, need to determine if `arresting_officers` is in *full-name-format* or *first-initial-format*

```
In [3]: is_uppercase <- function(chr) {
  res <- chr %in% LETTERS
  return(res)
}
```

```

has_full_names <- function(names_str) {
  char1 <- substr(names_str, 1, 1)
  char2 <- substr(names_str, 2, 2)

  res <- !(is_uppercase(char1) && is_uppercase(char2))
  return(res)
}

```

### 3.1.1 3.1.1 Counting the Names

- Want to correct count names regardless of format
- Update our count\_names() function

```

In [5]: count_names <- function(names_str) {
  names_str_trm <- str_trim(names_str)      # remove whitespace

  if (has_full_names(names_str_trm)) {
    split_char <- "/"
  } else {
    split_char <- ", "
  }

  name_vec <- unlist(str_split(names_str_trm, split_char))
  k <- length(name_vec)

  return(k)
}

```

### 3.1.2 3.1.2 Counting Officers (correctly)

- Note the sequence of function calls:

```
- count_officers() => count_names() => has_full_names() => is_uppercase()
```

```

In [6]: count_officers <- function(col) {

  n <- length(col)      # get the length of our input column
  cnts <- rep(0, n)     # allocate vector of zeros to populate with counts

  for (i in 1:n) {
    cnts[i] <- count_names(col[i])
  }
  return(cnts)
}

```

```
In [7]: arrests_df$officer_cnt <- count_officers(arrests_df$arresting_officers)
```

## 3.2 3.2 Add Violent Offense Flag

```
In [8]: # Write function to flag alleged violent crimes from the
# description of of the statute violation

is_violent_offense <- function(v) {

  violent_terms <- c("domestic-asslt", "assault", "battery", "murder")
  n_obs <- length(v)
  is_violent <- rep(FALSE, n_obs)

  # iterate over all statute descriptions
  for (i in 1:n_obs) {

    # iterate over the 4 terms associated with violence
    for (term in violent_terms) {
      if (!is.na(v[i]) && str_detect(tolower(v[i]), term)) {

        is_violent[i] <- TRUE
      }
    }
  }
  return(is_violent)
}
```

### 3.2.1 3.2.1 Test our Function (always!!)

```
In [9]: vio_vec <- c("DISORDERLY CONDUCT",
"RESISTING LEGAL OR ILLEGAL ARREST",
"DOMESTIC-SIMPLE ASSAULT/BATTERY",
"SIMPLE ASSAULT OR BATTERY")

is_violent_offense(vio_vec) # Should be: FALSE, FALSE, TRUE, TRUE
```

1. FALSE 2. FALSE 3. TRUE 4. TRUE

### 3.2.2 3.2.2 Create violent Column in arrests\_df

```
In [10]: arrests_df$violent <- is_violent_offense(arrests_df$statute_desc)
```

```
In [12]: head(arrests_df)
```

A data.frame: 6 CE 20

|   | arrest_date<br><chr>  | year<br><int> | month<br><int> | gender<br><chr> | race<br><chr> | ethnicity<br><chr> | year_of<br><int> |
|---|-----------------------|---------------|----------------|-----------------|---------------|--------------------|------------------|
| 1 | 2019-08-24T02:23:00.0 | 2019          | 8              | Male            | White         | NonHispanic        | 1981             |
| 2 | 2019-08-24T02:02:00.0 | 2019          | 8              |                 |               |                    | 1994             |
| 3 | 2019-08-24T02:02:00.0 | 2019          | 8              | Female          | Black         | NonHispanic        | 1984             |
| 4 | 2019-08-24T02:02:00.0 | 2019          | 8              | Female          | Black         | NonHispanic        | 1984             |
| 5 | 2019-08-24T02:02:00.0 | 2019          | 8              | Female          | Black         | Unknown            | 2001             |
| 6 | 2019-08-24T02:02:00.0 | 2019          | 8              | Female          | Black         | Unknown            | 2001             |

### 3.3 Aggregating to *Person-Level* DataFrame

- Use the `group_by()` and `summarise()` pattern from *dplyr* functions

```
In [14]: person_df <- arrests_df %>%  
  group_by(arrestee_id) %>%  
  summarise(  
    total_charges = n(),  
    num_uniq_arrests = length(unique(case_number)),  
    prop_violent = mean(violent),  
    mean_officer_cnt = mean(officer_cnt),  
    age = age[1],  
    gender = gender[1]  
  )  
  
`summarise()` ungrouping output (override with `.groups` argument)
```

```
In [15]: head(person_df)
```

|                 | arrestee_id<br><chr>    | total_charges<br><int> | num_uniq_arrests<br><int> | prop_violent<br><dbl> | mean_officer_cnt<br><dbl> |
|-----------------|-------------------------|------------------------|---------------------------|-----------------------|---------------------------|
| A tibble: 6 x 7 | pvd10005240635544439514 | 2                      | 1                         | 0                     | 1.000000                  |
|                 | pvd10007039892056892673 | 1                      | 1                         | 0                     | 1.000000                  |
|                 | pvd10015003399035869819 | 6                      | 2                         | 0                     | 4.000000                  |
|                 | pvd10015761183771579680 | 1                      | 1                         | 0                     | 2.000000                  |
|                 | pvd10016651127192901464 | 1                      | 1                         | 1                     | 2.000000                  |
|                 | pvd10028326204653807523 | 3                      | 3                         | 0                     | 1.666667                  |

## 4. Intro to *ggplot2*

- Operates on `data.frame` objects
- Map variables to aesthetics, and then display using “geom” (i.e., “geometric object”)
- Geom layers can be stacked over one another to add information

```
In [16]: ggplot(person_df, aes(x = age))    # does nothing...
```

#### 4.1 4.1 Plotting Histogram of age

```
In [17]: ggplot(person_df, aes(x = age)) +  
         geom_histogram()    # kinda boring...  
  
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

#### 4.1.1 4.1.1 Adding colour and fill to geom\_histogram()

```
In [19]: ggplot(person_df, aes(x = age)) +  
         geom_histogram(fill = "skyblue", colour = "lightblue", bins = 30)
```

## 4.2 4.2 Density Plot of age

```
In [21]: ggplot(person_df, aes(x = age)) +  
         geom_density(fill = "skyblue", colour = "lightblue")
```

#### 4.2.1 4.2.1 Adjusting alpha

```
In [22]: ggplot(person_df, aes(x = age)) +  
         geom_density(fill = "skyblue", colour = "lightblue", alpha = 0.5)
```



#### 4.2.2 4.2.2 Adding gender Variable Aesthetic

```
In [23]: ggplot(person_df, aes(x = age, y = stat(count), fill = gender, colour = gender)) +  
  geom_density(alpha = 0.4) +  
  xlab("Age of Person Arrested") +  
  ylab("Count")
```

### 4.3 4.3 Scatter Plot of age and total\_charges

```
In [24]: ggplot(person_df, aes(x = age, y = total_charges)) +  
         geom_point()
```

#### 4.3.1 4.3.1 Adjusting colour and alpha

```
In [25]: ggplot(person_df, aes(x = age, y = total_charges)) +  
         geom_point(colour = "skyblue", alpha = 0.6)
```

#### 4.3.2 4.3.2 Using geom\_jitter for Scatterplots

```
In [30]: ggplot(person_df, aes(x = age, y = total_charges)) +  
         geom_jitter(colour = "purple", alpha = 0.4)
```

#### 4.4 4.4 Plotting num\_uniq\_arrests and total\_charges with a stat\_smooth() Layer

```
In [31]: ggplot(person_df, aes(x = num_uniq_arrests, y = total_charges)) +  
  geom_jitter(colour = "violet", alpha = 0.4) +  
  stat_smooth(method = "lm", formula = y ~ x)
```

#### 4.5 4.5 Adding Third Variable to aes()

```
In [33]: ggplot(person_df, aes(x = num_uniq_arrests, y = total_charges, colour = prop_violent))  
         geom_jitter(alpha = 0.5) +  
         xlim(1, 6) +  
         ylim(1, 15)
```

Warning message:

Removed 2585 rows containing missing values (geom\_point).

