

sql_joins

November 23, 2020

Using SQL JOIN to Query Multiple Tables

1 1. JOIN Allows us to Expand our Query's Result Set

- If you'll recall, we used `leftjoin()` from *dplyr* in R
 - This function was inspired by SQL
 - Frequently hear the term “SQL-style join”

1.1 1.1 Some Terminology

- Joins are performed by using a given column that appears in both tables being joined
- Frequently, the column being used to join on is “primary key” or a “foreign key”
 - A primary key is the column that acts as the unique identifier for a given record (i.e., row)
 - A foreign key is a column in a table that is associated with a primary key in another table
- Caveat:
 - Both primary keys and foreign keys can be a combination of columns (ignore this for now)

1.2 1.2 Example JOIN

- Recall our `amzn.db` database has several tables
 - `orders`
 - `products`
 - `customer`
- Suppose we want to know what each customer paid for their order
- This will involve a query across `orders` and `products`, since the price information is in the `products` table

In [2]: `%load data/amzn.db`

1.2.1 1.2.1 Quick Refresher on SQL SELECT

```
In [3]: SELECT
        customer,      -- these are the columns (i.e., "fields") we want
        date,
        product_id
FROM
    orders              -- this is the table we pull from
;
```

```
Out [3]: +-----+-----+-----+
| customer | date       | product_id |
+-----+-----+-----+
| lee      | 2018-12-23 | 44         |
+-----+-----+-----+
| smith    | 2020-03-12 | 33         |
+-----+-----+-----+
| jones    | 2019-05-01 | 212        |
+-----+-----+-----+
| yang     | 2020-09-12 | 12         |
+-----+-----+-----+
| guerra   | 2020-08-03 | 12         |
+-----+-----+-----+
| diaz     | 2020-11-28 | 123        |
+-----+-----+-----+
| riley    | 2019-05-18 | 232        |
+-----+-----+-----+
| chan     | 2018-10-03 | 28         |
+-----+-----+-----+
```

1.2.2 1.2.2 Recall the SELECT * Idiom

- This will show you all rows and columns for a given table(s)
- Not usually advisable

```
In [6]: SELECT
        *
FROM
    products
;
```

```
Out [6]: +-----+-----+-----+-----+
| product_id | product   | price  | supplier_id |
+-----+-----+-----+-----+
| 44         | rake      | 19.99  | 65656       |
+-----+-----+-----+-----+
| 33         | shoe horn | 5.99   | 3434        |
+-----+-----+-----+-----+
| 22         | potato    | 0.99   | 65656       |
+-----+-----+-----+-----+
```

12	bike	123.5	232	
123	table	78.55	54545	
232	cup	2.5	4333	
28	ball	5.5	2323	
22	pencil	2.99	3232	
11	teapot	12.49	6565	
13	fork	1.99	86787	
14	shoelace	0.5	8787	
555	hammer	17.49	7878	
66	door	159.99	9889	

1.3 1.3 Motivating Example

- Suppose we want to obtain the total cost of every order
- Note that our orders table does not contain price information
- The orders table does have a product_id field
 - the product_id field can be used to link (i.e., “join”) to the product table

```
In [7]: SELECT
        customer,
        quantity,
        product,
        price
FROM
    orders
    JOIN products ON orders.product_id = products.product_id
;
```

```
Out[7]: +-----+-----+-----+-----+
| customer | quantity | product  | price |
+-----+-----+-----+-----+
| lee      | 1        | rake     | 19.99 |
+-----+-----+-----+-----+
| smith    | 1        | shoe horn | 5.99  |
+-----+-----+-----+-----+
| yang     | 1        | bike     | 123.5 |
+-----+-----+-----+-----+
```

guerra	2	bike	123.5	
+-----+	+-----+	+-----+	+-----+	+-----+
diaz	2	table	78.55	
+-----+	+-----+	+-----+	+-----+	+-----+
riley	4	cup	2.5	
+-----+	+-----+	+-----+	+-----+	+-----+
chan	1	ball	5.5	
+-----+	+-----+	+-----+	+-----+	+-----+

1.3.1 1.3.1 Using Table Aliases

- It is more idiomatic to alias our table names to something short (e.g., ord for orders)
- We can then use that in the SELECT section of our query

```
In [8]: SELECT
        ord.customer,
        ord.quantity,
        ord.product_id,
        pro.product,
        pro.price
FROM
        orders          AS ord
        JOIN products AS pro ON ord.product_id = pro.product_id
;
```

```
Out [8]: +-----+-----+-----+-----+-----+
| customer | quantity | product_id | product  | price |
+-----+-----+-----+-----+-----+
| lee      | 1        | 44         | rake     | 19.99 |
+-----+-----+-----+-----+-----+
| smith    | 1        | 33         | shoe horn | 5.99  |
+-----+-----+-----+-----+-----+
| yang     | 1        | 12         | bike     | 123.5 |
+-----+-----+-----+-----+-----+
| guerra   | 2        | 12         | bike     | 123.5 |
+-----+-----+-----+-----+-----+
| diaz     | 2        | 123        | table    | 78.55 |
+-----+-----+-----+-----+-----+
| riley    | 4        | 232        | cup      | 2.5   |
+-----+-----+-----+-----+-----+
| chan     | 1        | 28         | ball     | 5.5   |
+-----+-----+-----+-----+-----+
```

1.3.2 1.3.2 Doing Math in SELECT Section

```
In [9]: SELECT
        ord.customer,
        ord.quantity,
        ord.product_id,
```

```

        pro.product,
        pro.price,
        (pro.price * ord.quantity) AS total_cost
FROM
    orders          AS ord
    JOIN products AS pro ON ord.product_id = pro.product_id
;

```

Out [9]:

customer	quantity	product_id	product	price	total_cost
lee	1	44	rake	19.99	19.99
smith	1	33	shoe horn	5.99	5.99
yang	1	12	bike	123.5	123.5
guerra	2	12	bike	123.5	247.0
diaz	2	123	table	78.55	157.1
riley	4	232	cup	2.5	10.0
chan	1	28	ball	5.5	5.5

1.3.3 Filtering Result using WHERE

- Suppose we only care about order with 2 or more items

In [10]:

```

SELECT
    ord.customer,
    ord.quantity,
    ord.product_id,
    pro.product,
    pro.price,
    (pro.price * ord.quantity) AS total_cost
FROM
    orders          AS ord
    JOIN products AS pro ON ord.product_id = pro.product_id
WHERE
    ord.quantity > 1
;

```

Out [10]:

customer	quantity	product_id	product	price	total_cost
guerra	2	12	bike	123.5	247.0

diaz	2	123	table	78.55	157.1	
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+
riley	4	232	cup	2.5	10.0	
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+

1.3.4 Filterin on Multiple Criteria

- Recall that we can use the WHERE clause to filter according to any number of criteria
- Suppose we want either orders with more than 2 items *or* those with a *particular* item

```
In [12]: SELECT
        ord.customer,
        ord.quantity,
        ord.product_id,
        pro.product,
        pro.price,
        (pro.price * ord.quantity) AS total_cost
FROM
        orders          AS ord
        JOIN products AS pro ON ord.product_id = pro.product_id
WHERE
        ord.quantity > 1
        OR pro.product IN ('rake', 'bike')
;
```

```
Out[12]: +-----+-----+-----+-----+-----+-----+
| customer | quantity | product_id | product | price | total_cost |
+-----+-----+-----+-----+-----+-----+
| lee      | 1        | 44         | rake    | 19.99 | 19.99      |
+-----+-----+-----+-----+-----+-----+
| yang     | 1        | 12         | bike    | 123.5 | 123.5      |
+-----+-----+-----+-----+-----+-----+
| guerra   | 2        | 12         | bike    | 123.5 | 247.0      |
+-----+-----+-----+-----+-----+-----+
| diaz     | 2        | 123        | table   | 78.55 | 157.1      |
+-----+-----+-----+-----+-----+-----+
| riley    | 4        | 232        | cup     | 2.5   | 10.0       |
+-----+-----+-----+-----+-----+-----+
```